


# **Deep Generative Models (Fall 2024)**

**CS HUFS**


# Generative Adversarial Networks I

- Likelihood-Free Learning
- Two-Sample Hypothesis Test
- Generative Adversarial Networks (GAN)
- Training GAN
- Mode Collapse

# Recap



$$\begin{aligned} \mathbf{x}_i &\sim P_{\text{data}} \\ i &= 1, 2, \dots, n \end{aligned}$$



- Model families

- Autoregressive Models:  $p_{\theta}(\mathbf{x}) = \prod_{i=1}^n p_{\theta}(x_i | \mathbf{x}_{<i})$
- Variational Autoencoders:  $p_{\theta}(\mathbf{x}) = \int p_{\theta}(\mathbf{x}, \mathbf{z}) d\mathbf{z}$
- Normalizing Flow Models:  $p_{\mathbf{x}}(\mathbf{x}; \theta) = p_{\mathbf{z}}(\mathbf{f}_{\theta}^{-1}(\mathbf{x})) \left| \det \left( \frac{\partial \mathbf{f}_{\theta}^{-1}(\mathbf{x})}{\partial \mathbf{x}} \right) \right|$

- All the above families are trained by minimizing KL divergence  $D_{KL}(p_{\text{data}} || p_{\theta})$ , or equivalently maximizing likelihoods (or approximations)

# Why maximum likelihood?

$$\hat{\theta} = \underset{\theta}{\operatorname{argmax}} \sum_{i=1}^M \log p_{\theta}(\mathbf{x}_i), \quad \mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_M \sim p_{\text{data}}(\mathbf{x})$$

- **Optimal statistical efficiency.**

- Assume sufficient model capacity, such that there exists a unique  $\theta^* \in M$  that satisfies  $p_{\theta^*} = p_{\text{data}}$ .
- The convergence of  $\hat{\theta}$  to  $\theta^*$  when  $M \rightarrow \infty$  is the “fastest” among all statistical methods when using maximum likelihood training.

- **Higher likelihood = better lossless compression.**

- Is the likelihood a good indicator of the quality of samples generated by the model?

# Towards likelihood-free learning

- **Case 1:** Optimal generative model will give best **sample quality** and highest test **log-likelihood**
- For imperfect models, achieving high log-likelihoods might not always imply good sample quality, and vice-versa (Theis et al., 2016)

# Towards likelihood-free learning

- **Case 2:** Great test log-likelihoods, poor samples. E.g., For a discrete noise mixture model  $p_{\theta}(\mathbf{x}) = 0.01p_{\text{data}}(\mathbf{x}) + 0.99p_{\text{noise}}(\mathbf{x})$ 
  - 99% of the samples are just noise (most samples are poor)
  - Taking logs, we get a lower bound

$$\begin{aligned}\log p_{\theta}(\mathbf{x}) &= \log[0.01p_{\text{data}}(\mathbf{x}) + 0.99p_{\text{noise}}(\mathbf{x})] \\ &\geq \log 0.01p_{\text{data}}(\mathbf{x}) = \log p_{\text{data}}(\mathbf{x}) - \log 100\end{aligned}$$

- For expected log-likelihoods, we know that
  - Lower bound

$$E_{p_{\text{data}}}[\log p_{\theta}(\mathbf{x})] \geq E_{p_{\text{data}}}[\log p_{\text{data}}(\mathbf{x})] - \log 100$$

- Upper bound (via non-negativity of  $D_{\text{KL}}(p_{\text{data}} \parallel p_{\theta}) \geq 0$ )

$$E_{p_{\text{data}}}[\log p_{\text{data}}(\mathbf{x})] \geq E_{p_{\text{data}}}[\log p_{\theta}(\mathbf{x})]$$

- As we increase the dimension  $n$  of  $\mathbf{x} = (x_1, \dots, x_n)$ , absolute value of  $\log p_{\text{data}}(\mathbf{x}) = \sum_{i=1}^n \log p_{\theta}(x_i | \mathbf{x}_{<i})$  increases proportionally to  $n$  but  $\log 100$  remains constant. Hence, likelihoods are great  
 $E_{p_{\text{data}}}[\log p_{\theta}(\mathbf{x})] \approx E_{p_{\text{data}}}[\log p_{\text{data}}(\mathbf{x})]$  in very high dimensions


# Towards likelihood-free learning

- **Case 3:** Great samples, poor test log-likelihoods. E.g., Memorizing training set
  - Samples look exactly like the training set (cannot do better!)
  - Test set will have zero probability assigned (cannot do worse!)
- The above cases suggest that it might be useful to disentangle likelihoods and sample quality
- **Likelihood-free learning** consider alternative training objectives that do not depend directly on a likelihood function

# Recap



$$\mathbf{x}_i \sim P_{\text{data}} \\ i = 1, 2, \dots, n$$



## Model families

- Autoregressive Models:  $p_{\theta}(\mathbf{x}) = \prod_{i=1}^n p_{\theta}(x_i | \mathbf{x}_{<i})$
- Variational Autoencoders:  $p_{\theta}(\mathbf{x}) = \int p_{\theta}(\mathbf{x}, \mathbf{z}) d\mathbf{z}$
- Normalizing Flow Models:  $p_{\mathbf{x}}(\mathbf{x}; \theta) = p_{\mathbf{z}}(\mathbf{f}_{\theta}^{-1}(\mathbf{x})) \left| \det\left(\frac{\partial \mathbf{f}_{\theta}^{-1}(\mathbf{x})}{\partial \mathbf{x}}\right) \right|$
- All the above families are trained by minimizing KL divergence  $d_{KL}(p_{\text{data}} || p_{\theta})$ , or equivalently maximizing likelihoods (or approximations)
- Today: alternative choices for  $d(p_{\text{data}} || p_{\theta})$



# Comparing distributions via samples



$$S_1 = \{\mathbf{x} \sim P\}$$

vs.



$$S_2 = \{\mathbf{x} \sim Q\}$$

Given a finite set of samples from two distributions  $S_1 = \{\mathbf{x} \sim P\}$  and  $S_2 = \{\mathbf{x} \sim Q\}$ , how can we tell if these samples are from the same distribution? (i.e.,  $P = Q$ )


# Two-sample tests

- Given  $S_1 = \{\mathbf{x} \sim P\}$  and  $S_2 = \{\mathbf{x} \sim Q\}$ , a **two-sample test** considers the following hypotheses
  - Null hypothesis  $H_0: P = Q$
  - Alternative hypothesis  $H_1: P \neq Q$
- Test statistic  $T$  compares  $S_1$  and  $S_2$ . For example: difference in means, variances of the two sets of samples
  - $$T(S_1, S_2) = \left| \frac{1}{|S_1|} \sum_{\mathbf{x} \in S_1} \mathbf{x} - \frac{1}{|S_2|} \sum_{\mathbf{x} \in S_2} \mathbf{x} \right|$$
- If  $T$  is larger than a threshold  $\alpha$ , then reject  $H_0$  otherwise we say  $H_0$  is consistent with observation.
- Key observation:** Test statistic is **likelihood-free** since it does not involve the densities  $P$  or  $Q$  (only samples)

# Generative modeling and two-sample tests



$x_i \sim P_{\text{data}}$   
 $i = 1, 2, \dots, n$




**Model family**

- A priori we assume direct access to  $S_1 = D = \{\mathbf{x} \sim p_{\text{data}}\}$
- In addition, we have a model distribution  $p_{\theta}$
- Assume that the model distribution permits efficient sampling (e.g., directed models). Let  $S_2 = \{\mathbf{x} \sim p_{\theta}\}$
- **Alternative notion of distance between distributions:** Train the generative model to minimize a two-sample test objective between  $S_1$  and  $S_2$


# Two-Sample Test via a Discriminator

- Finding a good two-sample test objective in high dimensions is hard



- In the generative model setup, we know that  $S_1$  and  $S_2$  come from different distributions  $p_{\text{data}}$  and  $p_{\theta}$  respectively
- Key idea:** Learn a statistic to automatically identify in what way the two sets of samples  $S_1$  and  $S_2$  differ from each other
- How? Train a classifier (called a discriminator)!

# Two-Sample Test via a Discriminator



## Two-Sample Test via a Discriminator

- Any binary classifier  $D_\phi$  (e.g., neural network) which tries to distinguish “real” ( $y = 1$ ) samples from the dataset and “fake” ( $y = 0$ ) samples generated from the model
- Test statistic: -loss of the classifier. Low loss, real and fake samples are easy to distinguish (different). High loss, real and fake samples are hard to distinguish (similar).
- Goal: Maximize the two-sample test statistic (in support of the alternative hypothesis  $p_{\text{data}} \neq p_\theta$ ), or equivalently minimize classification loss

# Two-Sample Test via a Discriminator

- Training objective for discriminator:

$$\begin{aligned}\max_{D_\phi} V(p_\theta, D_\phi) &= E_{\mathbf{x} \sim p_{\text{data}}}[\log D_\phi(\mathbf{x})] + E_{\mathbf{x} \sim p_\theta}[\log(1 - D_\phi(\mathbf{x}))] \\ &\approx \sum_{\mathbf{x} \in S_1} \log D_\phi(\mathbf{x}) + \sum_{\mathbf{x} \in S_2} [\log(1 - D_\phi(\mathbf{x}))]\end{aligned}$$


- For a fixed generative model  $p_\theta$ , the discriminator  $D_\phi$  is performing binary classification with the cross entropy objective
  - Assign probability 1 to **true data** points  $\mathbf{x} \sim p_{\text{data}}$  (in set  $S_1$ )
  - Assign probability 0 to **fake** samples  $\mathbf{x} \sim p_\theta$  (in set  $S_2$ )
- Optimal discriminator

$$D_\theta^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_\theta(\mathbf{x})}$$

- Sanity check: if  $p_\theta = p_{\text{data}}$ , classifier cannot do better than chance ( $D_\theta^*(\mathbf{x}) = 1/2$ )

# Generative Adversarial Networks

- A two player minimax game between a **generator** and a **discriminator**



- **Generator**

- Directed, latent variable model with a deterministic mapping between  $\mathbf{z}$  and  $\mathbf{x}$  given by  $G_\theta$ 
  - Sample  $\mathbf{z} \sim p(\mathbf{z})$ , where  $p(\mathbf{z})$  is a simple prior, e.g. Gaussian
  - Set  $\mathbf{x} = G_\theta(\mathbf{z})$
- Similar to a flow model, but mapping  $G_\theta$  need not be invertible
- Distribution over  $p_\theta(\mathbf{x})$  over  $\mathbf{x}$  is implicitly defined (no likelihood!)
- Minimizes a two-sample test objective (in support of the null hypothesis  $p_{\text{data}} = p_\theta$ )

# Example of GAN objective

- **Training objective for generator:**

$$\min_G \max_D V(G, D) = E_{\mathbf{x} \sim p_{\text{data}}} [\log D(\mathbf{x})] + E_{\mathbf{x} \sim p_G} [\log(1 - D(\mathbf{x}))]$$

- For the optimal discriminator  $D_G^*(\cdot)$ , we have

$$\begin{aligned} & V(G, D_G^*(\mathbf{x})) \\ &= E_{\mathbf{x} \sim p_{\text{data}}} \left[ \log \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})} \right] + E_{\mathbf{x} \sim p_G} \left[ \log \frac{p_G(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})} \right] \\ &= E_{\mathbf{x} \sim p_{\text{data}}} \left[ \log \frac{p_{\text{data}}(\mathbf{x})}{\frac{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})}{2}} \right] + E_{\mathbf{x} \sim p_G} \left[ \log \frac{p_G(\mathbf{x})}{\frac{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})}{2}} \right] - \log 4 \\ &= \underbrace{D_{KL} \left[ p_{\text{data}}, \frac{p_{\text{data}} + p_G}{2} \right] + D_{KL} \left[ p_G, \frac{p_{\text{data}} + p_G}{2} \right]}_{2 \times \text{Jensen-Shannon Divergence (JSD)}} - \log 4 \\ &= 2D_{\text{JSD}}[p_{\text{data}}, p_G] - \log 4 \end{aligned}$$



# Jenson-Shannon Divergence

- Also called as the symmetric KL divergence

$$D_{JSD}[p, q] = \frac{1}{2} \left( D_{KL} \left[ p, \frac{p+q}{2} \right] + D_{KL} \left[ q, \frac{p+q}{2} \right] \right)$$

- Properties

- $D_{JSD}[p, q] \geq 0$
- $D_{JSD}[p, q] = 0$  iff  $p = q$
- $D_{JSD}[p, q] = D_{JSD}[q, p]$
- $\sqrt{D_{JSD}[p, q]}$  satisfies triangle inequality  $\rightarrow$  Jenson-Shannon Distance

- Optimal generator for the JSD/Negative Cross Entropy GAN

$$p_G = p_{\text{data}}$$


- For the optimal discriminator  $D_{G^*}^*(\cdot)$  and generator  $G^*(\cdot)$ , we have

$$V(G^*, D_{G^*}^*(\mathbf{x})) = -\log 4$$

# Recap of GANs



$$\begin{aligned} \mathbf{x}_i &\sim P_{\text{data}} \\ i &= 1, 2, \dots, n \end{aligned}$$



Model family

- Choose  $d(p_{\text{data}}, p_{\theta})$  to be a two-sample test statistic
  - Learn the statistic by training a classifier (discriminator)
  - Under ideal conditions, equivalent to choosing  $d(p_{\text{data}}, p_{\theta})$  to be  $D_{\text{JSD}}[p_{\text{data}}, p_{\theta}]$
- Pros:
  - Loss only requires samples from  $p_{\theta}$ . No likelihood needed!
  - Lots of flexibility for the neural network architecture, any  $G_{\theta}$  defines a valid sampling procedure
  - Fast sampling (single forward pass)
- Cons: very difficult to train in practice

# The GAN training algorithm

- Sample minibatch of  $m$  training points  $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(m)}$  from  $D$
- Sample minibatch of  $m$  noise vectors  $\mathbf{z}^{(1)}, \mathbf{z}^{(2)}, \dots, \mathbf{z}^{(m)}$  from  $p_z$
- Update the discriminator parameters  $\phi$  by stochastic gradient **ascent**

$$\nabla_{\phi} V(G_{\theta}, D_{\phi}) = \frac{1}{m} \nabla_{\phi} \sum_{i=1}^m [\log D_{\phi}(\mathbf{x}^{(i)}) + \log(1 - D_{\phi}(G_{\theta}(\mathbf{z}^{(i)})))]$$

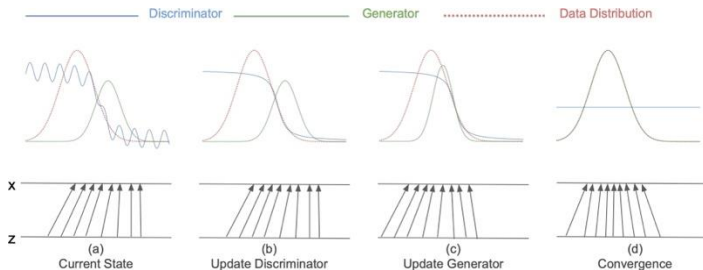
- Update the generator parameters  $\theta$  by stochastic gradient **descent**

$$\nabla_{\theta} V(G_{\theta}, D_{\phi}) = \frac{1}{m} \nabla_{\theta} \sum_{i=1}^m \log(1 - D_{\phi}(G_{\theta}(\mathbf{z}^{(i)})))$$


- Repeat for fixed number of epochs

# Alternating optimization in GANs

$$\min_{\theta} \max_{\phi} V(G_{\theta}, D_{\phi}) = E_{\mathbf{x} \sim p_{\text{data}}} [\log D_{\phi}(\mathbf{x})] + E_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D_{\phi}(G_{\theta}(\mathbf{z})))]$$



Which one is real?



Source: Karras et al., 2018; The New York Times

Both images are generated via GANs!

# Frontiers in GAN research



2014



2015



2016



2017



2018


- GANs have been successfully applied to several domains and tasks
- However, working with GANs can be very challenging in practice
  - Unstable optimization
  - Mode collapse
  - Evaluation
- Bag of tricks needed to train GANs successfully

---

Image Source: Ian Goodfellow. Samples from Goodfellow et al., 2014, Radford et al., 2015, Liu et al., 2016, Karras et al., 2017, Karras et al., 2018

# Optimization challenges

- **Theorem (informal):** If the generator updates are made in function space and discriminator is optimal at every step, then the generator is guaranteed to converge to the data distribution
- **Unrealistic assumptions!**
- In practice, the generator and discriminator loss keeps oscillating during GAN training



Source: Mirantha Jayathilaka

- No robust stopping criteria in practice (unlike MLE)

# Mode Collapse

- GANs are notorious for suffering from **mode collapse**
- Intuitively, this refers to the phenomena where the generator of a GAN collapses to one or few samples (dubbed as “modes”)




Arjovsky et al., 2017

“generates similar images, not diverse”




# Mode Collapse




Target


- True distribution is a mixture of Gaussians




Step 0



Step 5k



Step 10k



Step 15k




Step 20k

Source: Metz et al., 2017

- The generator distribution keeps oscillating between different modes


# Mode Collapse



Source: Metz et al., 2017

- Fixes to mode collapse are mostly empirically driven: alternative architectures, alternative GAN loss, adding regularization terms, etc.
- <https://github.com/soumith/ganhacks>  
How to Train a GAN? Tips and tricks to make GANs work by Soumith Chintala

# Beauty lies in the eyes of the discriminator



Source: Robbie Barrat, Obvious

GAN generated art auctioned at Christie's.

**Expected Price:** \$7,000 – \$10,000

**Price Sold:** \$432,500

Or was it? On Twitter, a 19-year-old programmer and artist named Robbie Barrat has accused the company of using open source code which he created.



# GAN Exercise 1

DATA

Dataset  
MNIST

Model - discriminator  
Convolutional (disc)

Model - generator  
Convolutional (gen)

Hyperparameters  
Learning Rate - discriminator  
0.01  
Optimizer - discriminator  
sgd  
Learning Rate - generator  
0.01  
Beta1 - generator  
0.9  
Beta2 - generator  
0.999  
Optimizer - generator  
adam  
Batch Size  
15  
Normalization  
[-1, 1]  
Statistics  
Examples 65000  
Input shape [28,28,1]  
Label shape [10]

DISCRIMINATOR

Input Image

[28,28,1]

Op type  
Convolution

Field sz... 5 Stride 1 Zero p... 2 Output d... 8

[28,28,8]

X

Op type  
ReLU

[28,28,8]

X

Op type  
Max pool

Field sz... 2 Stride 2 Zero p... 0

[14,14,8]

X

Op type  
Convolution

Field sz... 5 Stride 1 Zero p... 2 Output d... 16

[14,14,16]

GENERATOR

Generator Random Vector

[100]

Op type  
Fully connected

Hidden ... 784

[784]

X

Op type  
ReLU

[784]

X

Op type  
Reshape

Shape (comma separated)  
28, 28, 1

[28,28,1]

X

Op type  
Convolution

Field sz... 3 Stride 1 Zero p... 1 Output d... 10

[28,28,10]

REAL IMAGES

Inferences/sec: 138  
Inference duration: 7.43ms

1

99.6%

0

0.4%

6

99.9%

0

0.1%

6

98.6%

0

1.4%

9

98.9%

0

1.1%

5

99.9%

0

0.1%

5

99.8%

0

0.2%

GENERATED IMAGES

Generations/sec: 138  
Generation duration: 7.43ms

0

97.9%

1

2.1%

0

80.7%

1

19.3%

0

99.3%

1

0.7%

0

98.0%

1

2.0%

0

96.6%

1

3.4%

1

64.0%

0

36.0%

TRAIN STATS

Examples/sec: 182  
Examples trained: 30000  
Total time: 286.8 sec.

Discriminator Cost

2.0  
1.5  
1.0  
0.5  
0

0 500 1000 1500

Generator Cost

10  
8  
6  
4  
2  
0

0 500 1000 1500

Examples/sec


1500  
1000  
500  
0

0 500 1000 1500

# GAN Exercise 2

## GAN Lab

Data Distribution



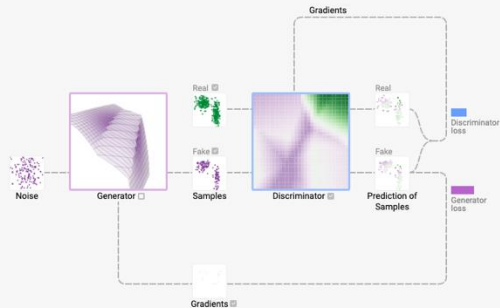
☒ Use pre-trained model




Epoch

003,298

### MODEL OVERVIEW GRAPH



### LAYERED DISTRIBUTIONS



Each dot is a 2D data sample: **real samples**, **fake samples**.


Background colors of grid cells represent **discriminator's** classifications. Samples in **green regions** are likely to be real; those in **purple regions** likely fake.

**Manifold** represents **generator's** transformation results from noise space. Opacity encodes density: darker purple means more samples in smaller area.

Pink lines from fake samples represent **gradients** for generator.

✓ This sample needs to move upper right to decrease generator's loss.

### METRICS



<https://poloclub.github.io/ganlab/>